

WELCOME

The Disruptor - High performance messaging with Java

Raffael Schmid

April 28, 2012

BASEL BERN LAUSANNE ZÜRICH DÜSSELDORF FRANKFURT A.M. FREIBURG I.BR. HAMBURG MÜNCHEN STUTTGART WIEN

1

2011 © Trivadis

The Disruptor - High performance messaging with Java
April 28, 2012

trivadis
makes IT easier. ■ ■ ■

AGENDA

1. About LMAX
2. How to tackle 6'000'000 transactions/second
3. The Disruptor

Where does the Disruptor come from?

It was initially designed to setup a trading platform called LMAX.

LMAX - platform

multiple devices

← accessible via

financial trading platform

July 2010

← since

LMAX platform

← is a

non-institutional clients

← accessible for

journaling

← needs

loss of data

← no

Financial Services Authority

← supervised

← needs to be

fast

6 mio. trx/sec

← means

← depend on

each other

replication

← needs

interruption

speed

← fundamental

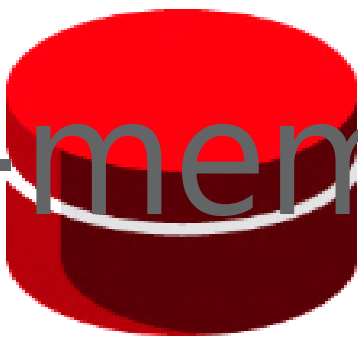
← the goal

the worlds fastest exchange



How NOT to take 600'000'000
transactions per second?

in-memory

A 3D red cylinder with a white horizontal line around its middle, positioned behind the text 'in-memory'.

How to tackle 6 mio. transactions / second

Software Transactional Memory

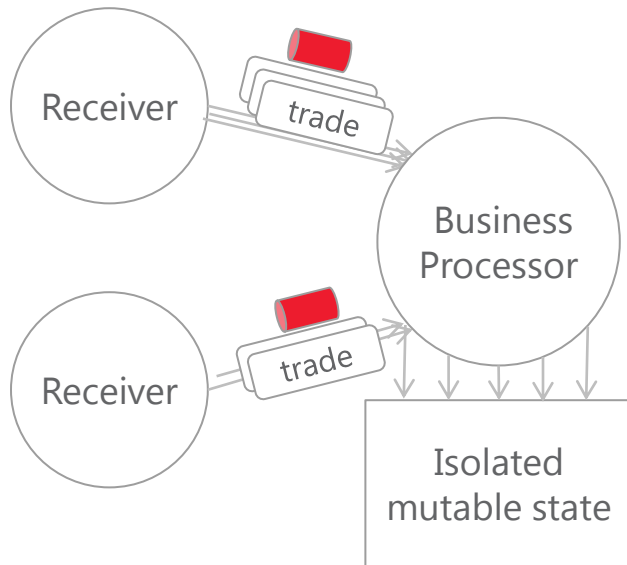
```
Ref<Integer> ref = new Ref<Integer>(0);  
Atomic<Integer> atomic = new  
    Atomic<Integer>() {  
    public Integer atomically() {  
        final int inc = ref.get() + 1;  
        ref.set(inc);  
        return inc;  
    }  
};  
atomic.execute();
```

- Turns Java heap into a transactional data set
- **A**tomicity
- **C**onsistency
- **I**solation: Serializable

 **does not scale**

How to tackle 6 mio. transactions / second

Actors

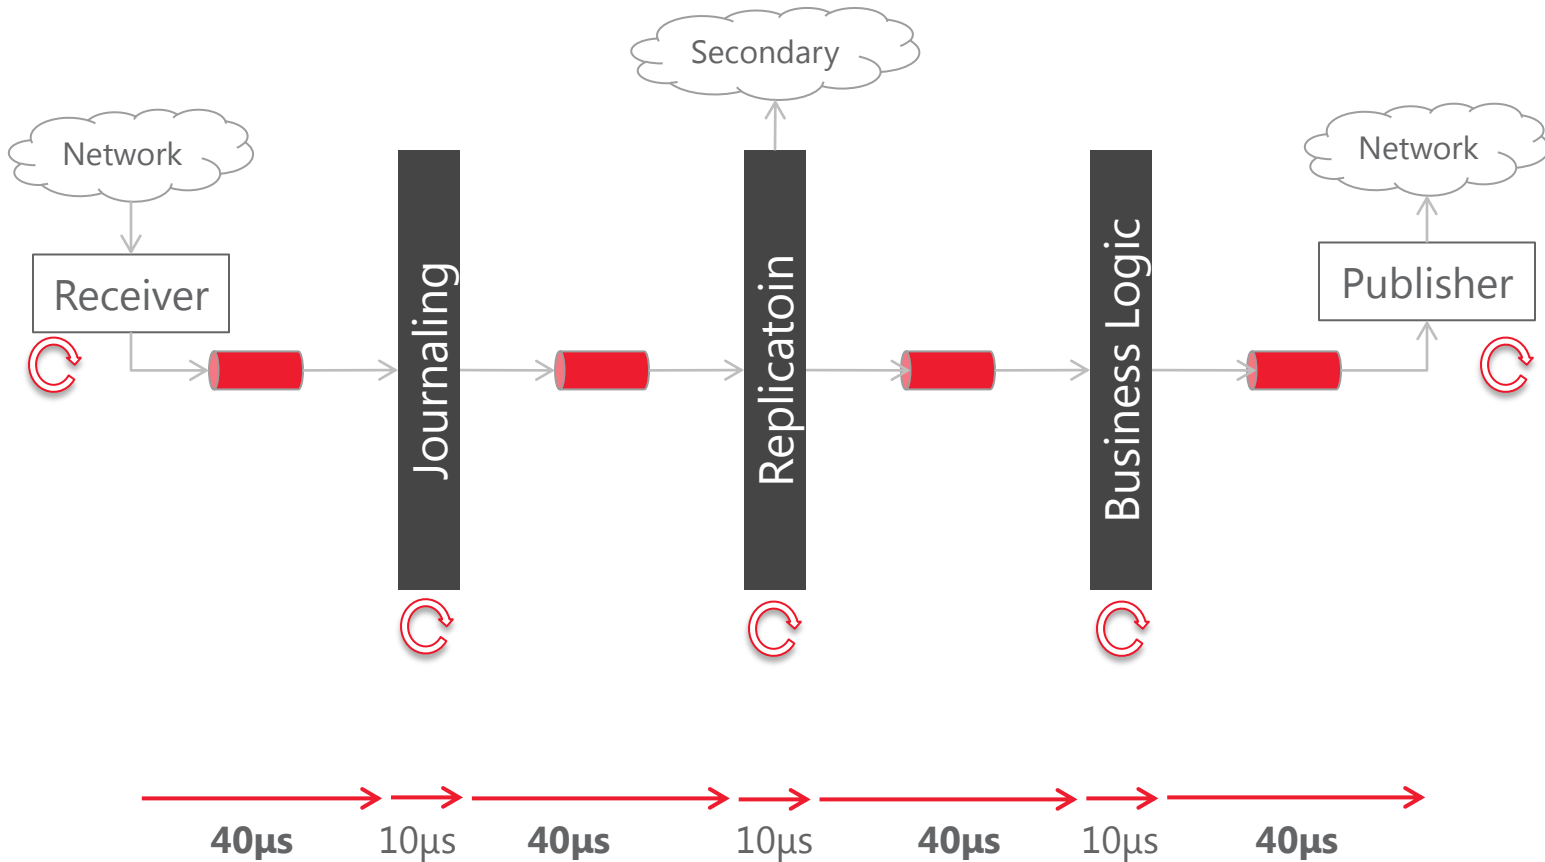


- Message passing between threads
- Queueing of messages
- Single-threaded BusinessProcessor

- **heavy contention on queues**
- **no guarantee for the sequence of events**
- **low predictable latency**

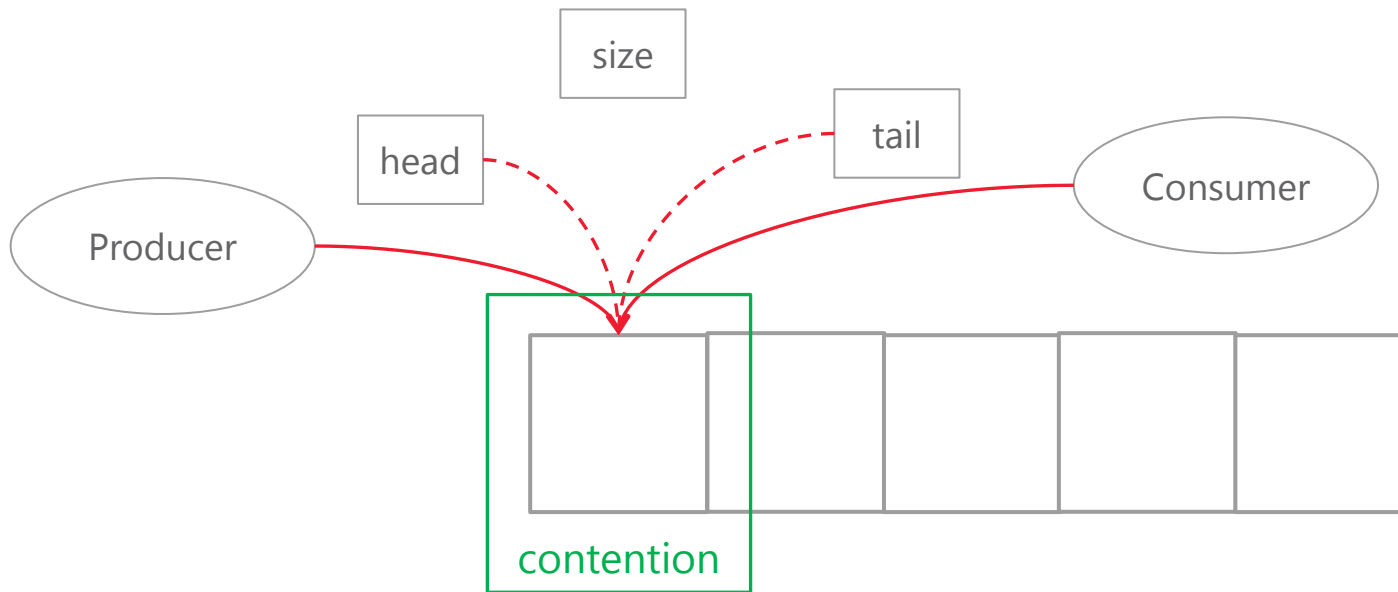
How to tackle 6 mio. transactions / second

Staged Event Driven Architecture



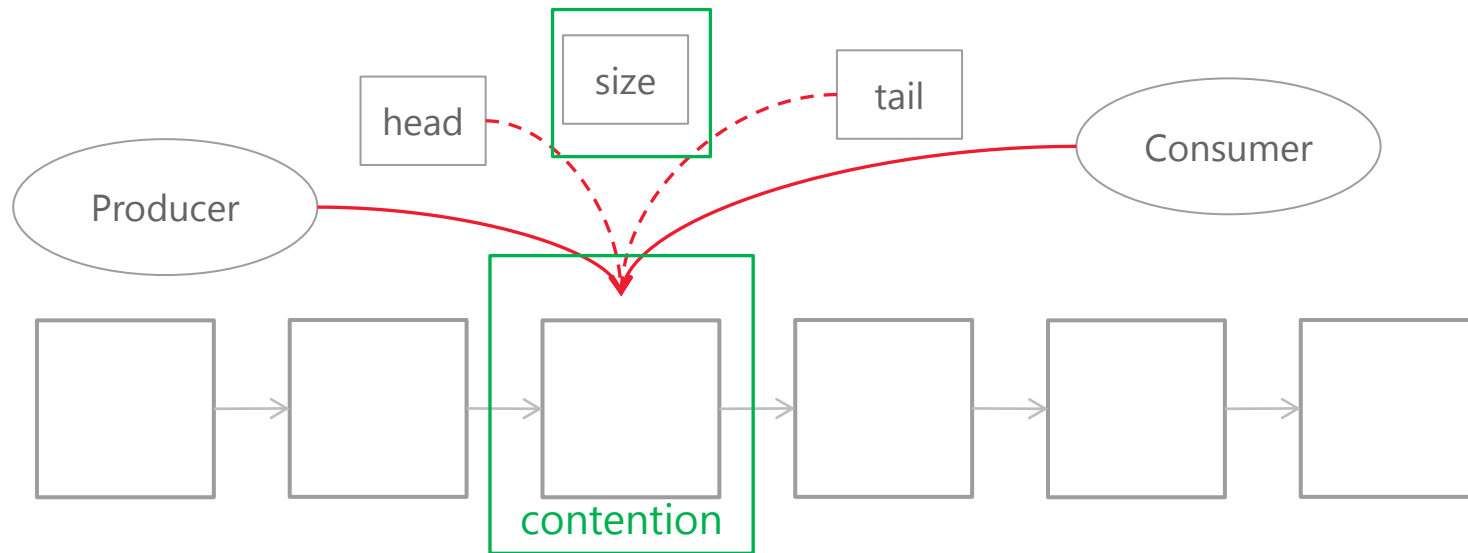
Why queues suck – Array backed

- Queues are never in steady state → **contention** on last or first element
- Head, tail may occupy same cache line → false sharing



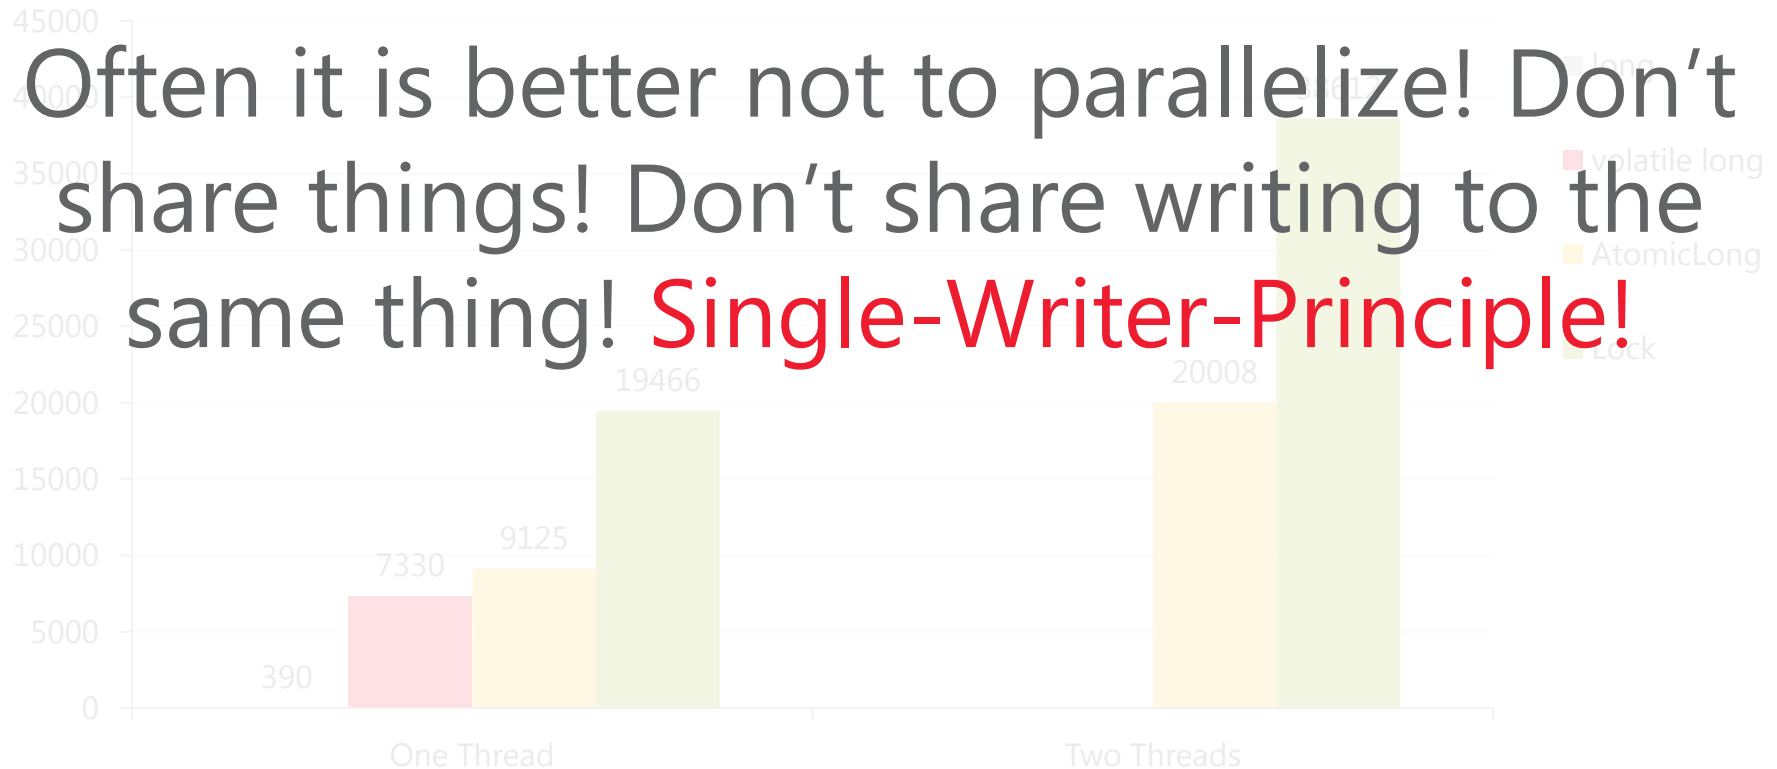
Why queues suck – LinkedList

- Not contiguous → do not support striding (prefetching)
- Needs to be bounded → contention on the size variable



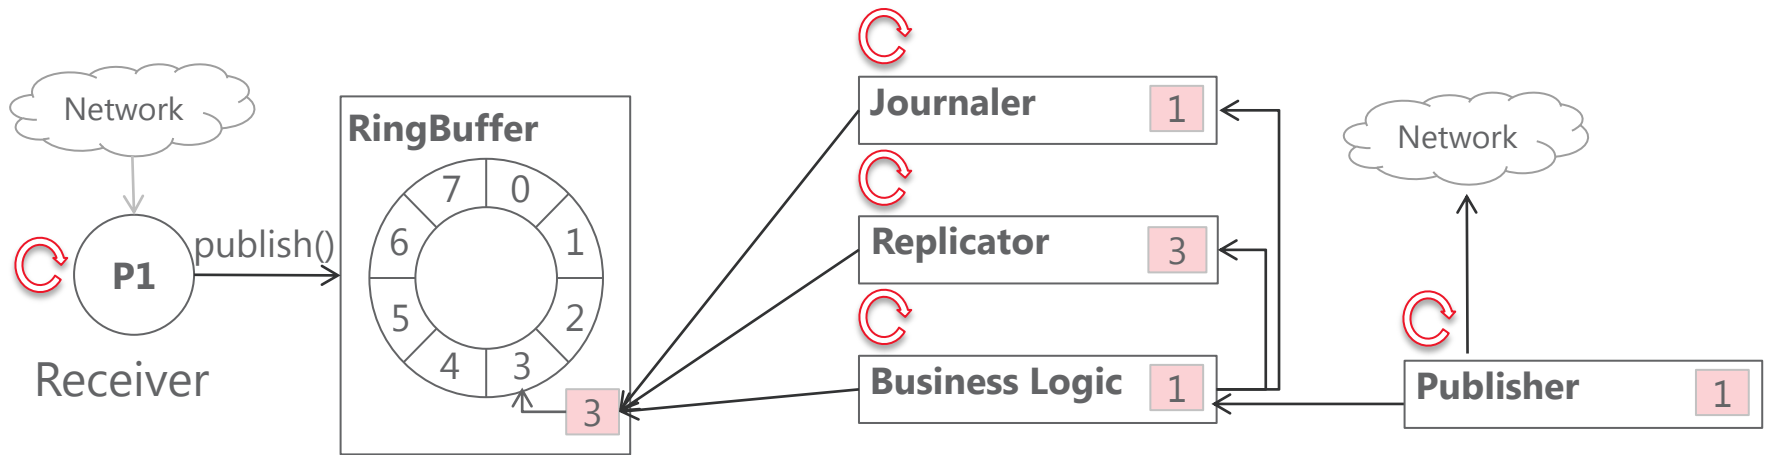
Costs of contention

Incrementing a counter 500 000 000 times.

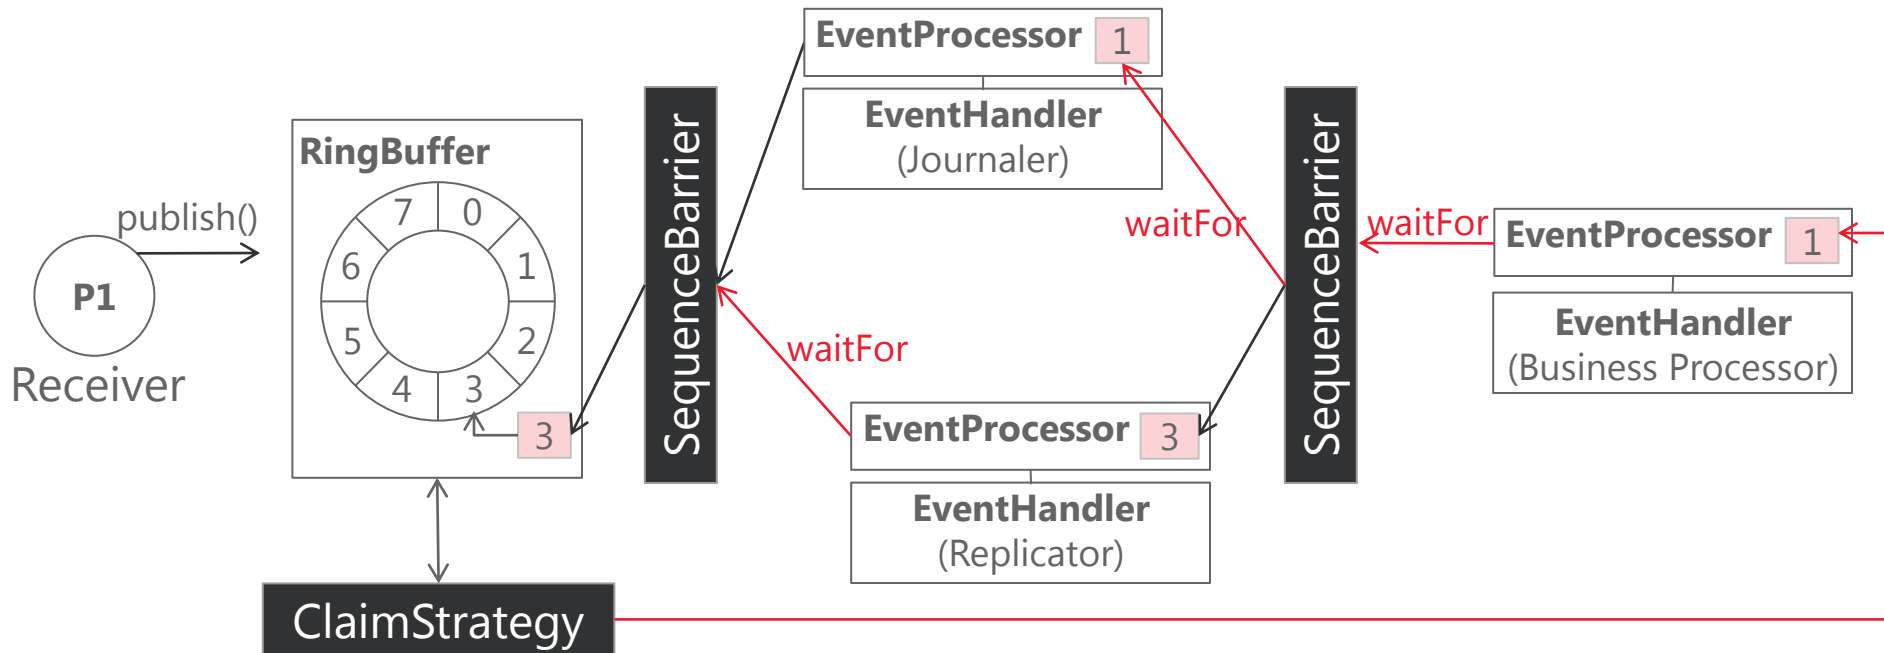


Results in milliseconds, run on a Intel Nehalem 2.93GHz

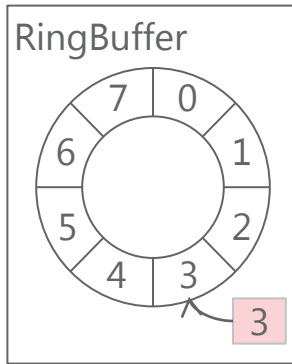
The Disruptor – contention-free design



The Disruptor – contention-free design



The Disruptor – RingBuffer

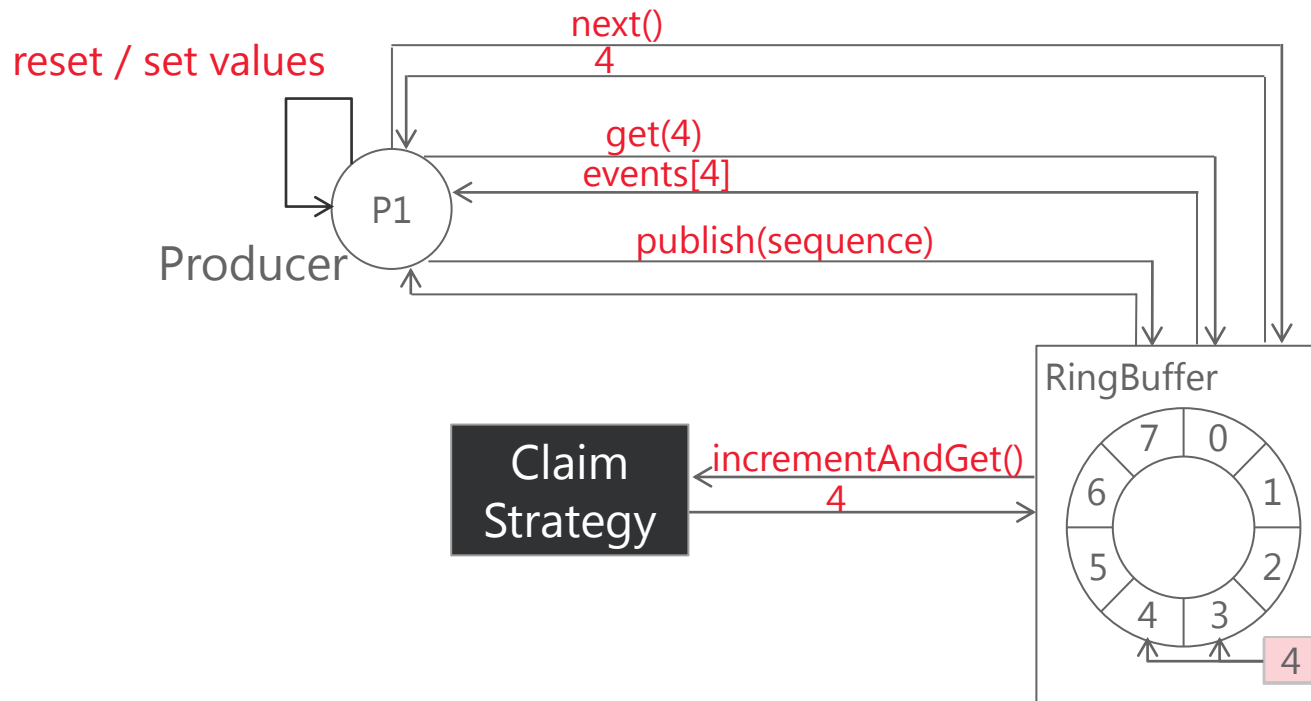


- **Garbage friendly**
- Based on an Array
 - Linear memory access → **prefetching**
 - Multiple entries per cache line
- Reliable messaging
- **Only** the thread writing to the RingBuffer is updating the cursor (Single-Writer-Principle)

**Entries created once
per slot**

**Only producers
update the cursor**

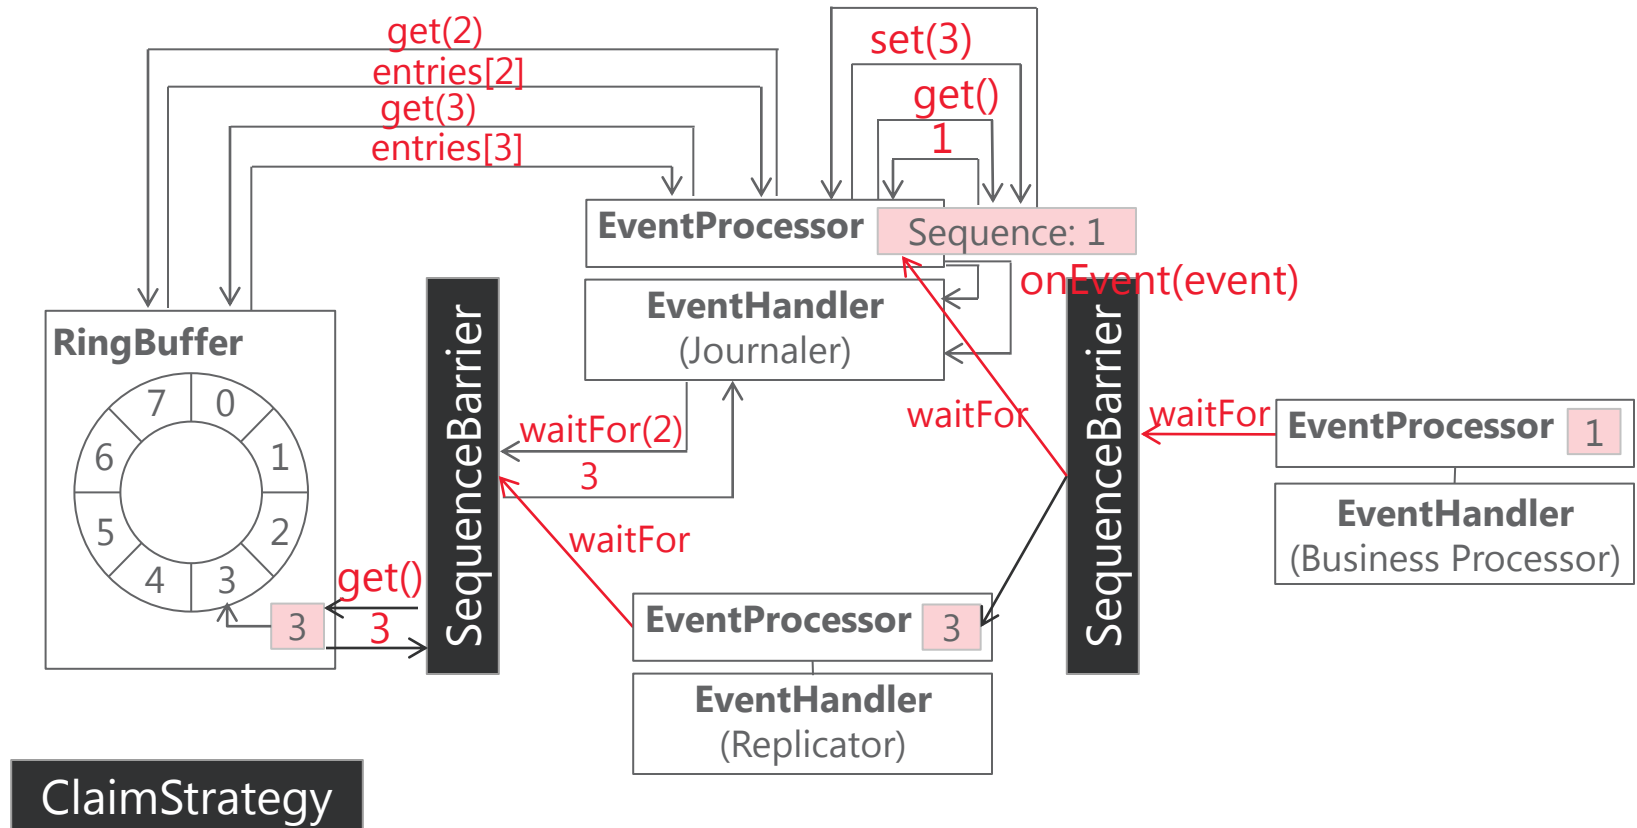
The Disruptor – writing to the RingBuffer



Two-Phase-Commit:
claim, publish

ClaimStrategy
inhibits wrapping

The Disruptor – Reading from the RingBuffer



Fields on a event should be written by only one EventProcessor!

Setup the Disruptor

```
RingBuffer<TradeEvent> buffer = new RingBuffer<TradeEvent>(
    eventFactory,
    claimStrategy,
    waitStrategy);
```

```
SequenceBarrier ringBarrier = buffer.newBarrier();
```

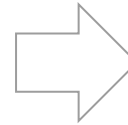
```
BatchEventProcessor<TradeEvent> replication = new BatchEventProcessor<TradeEvent>(
    buffer,
    ringBarrier,
    new ReplicationHandler());
BatchEventProcessor<TradeEvent> journaling = new BatchEventProcessor<TradeEvent>(
    buffer,
    ringBarrier,
    new JournalingHandler());
```

```
SequenceBarrier journalingReplicationBarrier = buffer.newBarrier(
    replProcessor.getSequence(),
    journProcessor.getSequence());
```

Publish messages

```
BatchEventProcessor<TradeEvent> businessLog = new BatchEventProcessor<TradeEvent>(
    buffer,
    journalingReplicationBarrier,
    blHandler);
buffer.setGatingSequences(businessLog.getSequence());
```

```
long sequence = buffer.next();
buffer.get(sequence).setValue(i);
buffer.publish(sequence);
```



2-phase commit

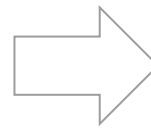
Setup the Disruptor using the DSL

```
Disruptor<TradeEvent> disruptor = new Disruptor<TradeEvent>(
    TradeEvent.EVENT_FACTORY,
    EXECUTOR,
    claimStrategy,
    waitStrategy);
```

```
disruptor.handleEventsWith( new ReplicationHandler(), new JournalingHandler())
    .then( new BusinessLogicHandler());

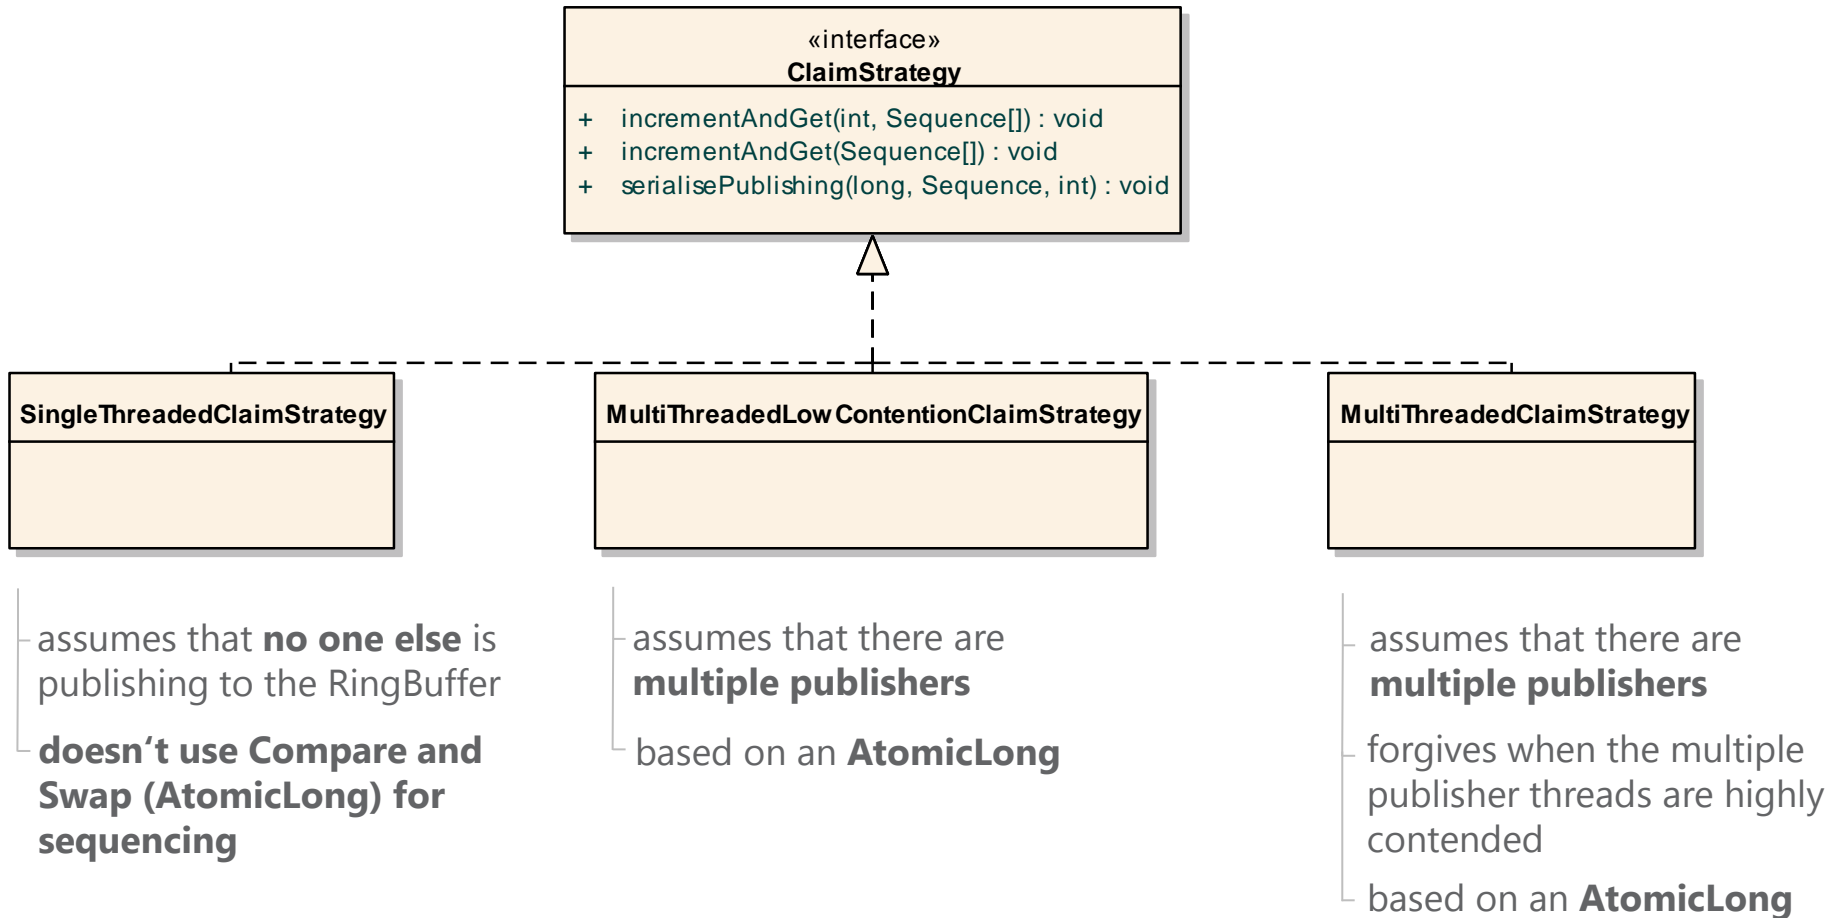
disruptor.start();
```

```
disruptor.publishEvent(eventTranslator);
```

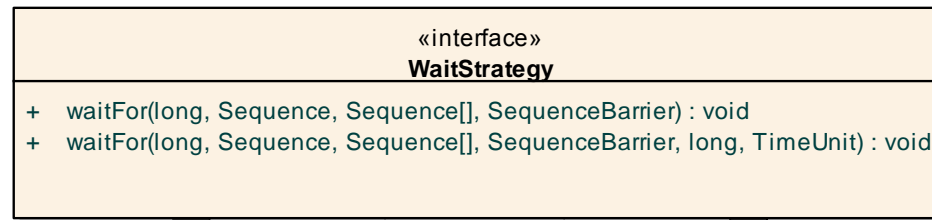


implicit Two-Phase-Commit

Which ClaimStrategy should I use?

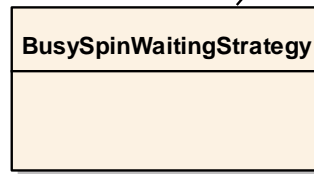


Decisions to be made - WaitStrategy



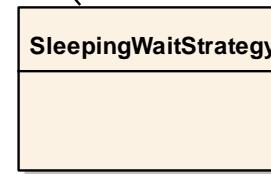
Puts current thread to state waiting until requested entry gets available

+ CPU resource
- Throughput



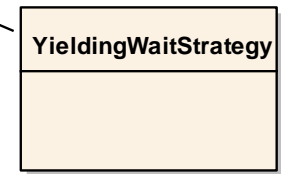
Spins until requested entry gets available

- CPU resource
+ Throughput



spin (100 times) → yield (100 times) → timed waiting

= CPU resource
= Throughput

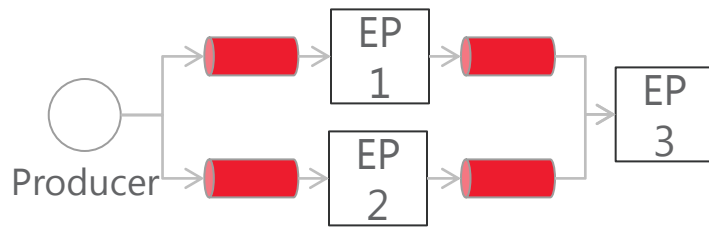


yields current thread

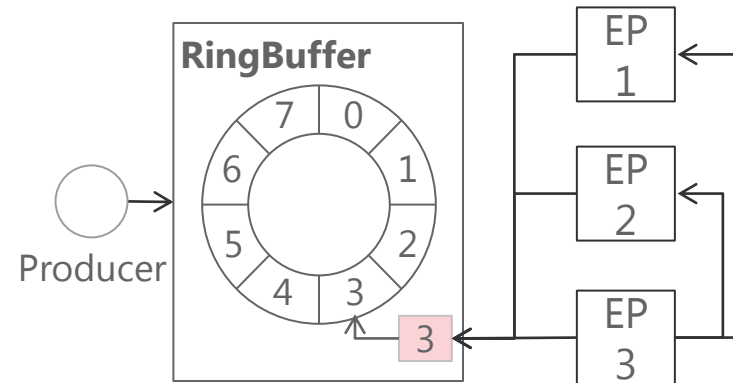
= CPU resource
= Throughput

Performance - LinkedBlockingQueue vs Disruptor

LinkedBlockingQueue

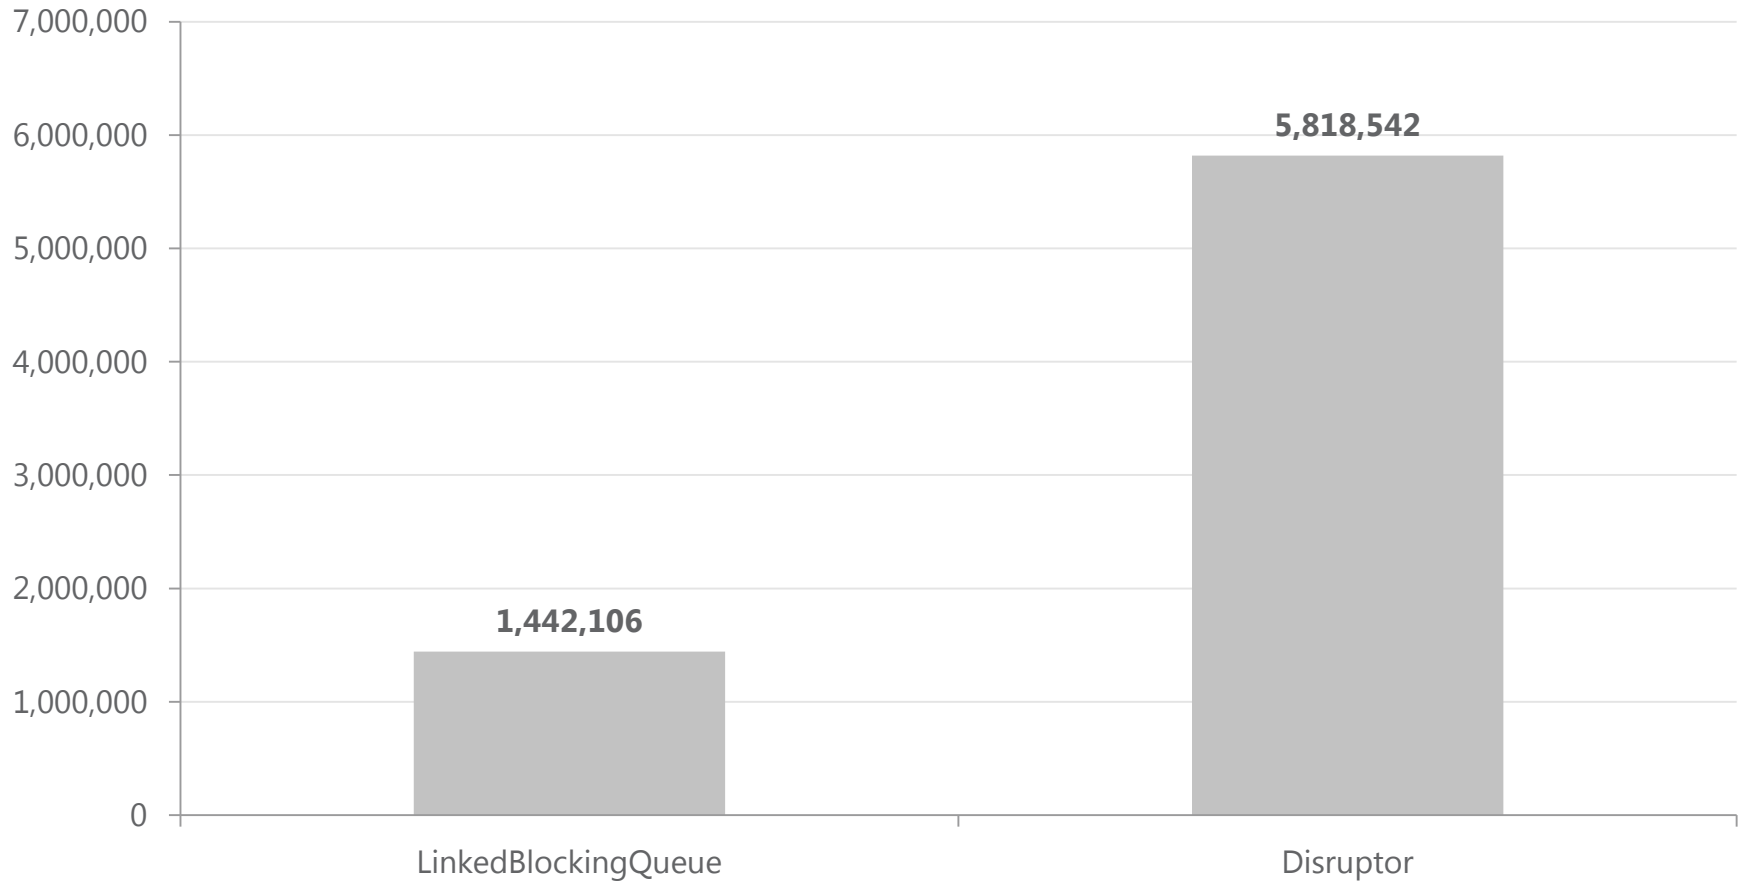


Disruptor



Performance - Throughput

Throughput



Why is it so fast

1. No Locks
2. minimized use of memory barriers
3. One single place of contention: if multiple producers write to the RingBuffer
4. Multiple producers, consumers share the same data structure
5. Recycling of events → less garbage → less garbage collection pauses
6. tracking sequences at each place + cache line padding → no false sharing → no unexpected contention

Sources

■ Blogs

- Martin Thompson
<http://mechanical-sympathy.blogspot.com>
- Michael Barker
<http://bad-concurrency.blogspot.com>
- Trisha Gee
<http://mechanitis.blogspot.com>

■ Articles

- Technical Paper
<http://disruptor.googlecode.com/files/Disruptor-1.0.pdf>
- Java Best Practices – Queue battle and the Linked ConcurrentHashMap
<http://www.javacodegeeks.com/2010/09/java-best-practices-queue-battle-and.html>

■ Presentations

- Understanding the Disruptor: A Beginners Guide to Hardcore Concurrency
<http://www.youtube.com/watch?v=DCdGlxBbKU4>

Questions



THANK YOU.

Trivadis AG

Raffael Schmid

Europa-Strasse 5
8152 Glattbrugg

Mobile +41 79 699 70 09
Tel. +41 44 808 70 20
Fax +41 44 808 70 21

raffael.schmid@trivadis.com
www.trivadis.com

BASEL BERN LAUSANNE ZÜRICH DÜSSELDORF FRANKFURT A.M. FREIBURG I.BR. HAMBURG MÜNCHEN STUTTGART WIEN